
title: "Final Dissertation code "

output: html_document

date: "2025-08-12"

```{r}

library(tidyverse)

library(lubridate)

library(tsibble)

library(feasts)

library(slider)

library(GGally)

library(anomalize)

library(vip)

library(dbSCAN)

theme\_set(theme\_minimal(base\_size = 12))

holiday\_dates <- as.Date(c(

"2010-11-26","2010-12-24",

"2011-11-25","2011-12-23",

"2012-11-23","2012-12-21"

))

```

01 Load & Clean

```{r }

```
train <- read_csv("train.csv", show_col_types = FALSE)
features <- read_csv("features.csv", show_col_types = FALSE)
stores <- read_csv("stores.csv", show_col_types = FALSE)

merged_data <- train %>%
 mutate(Date = as.Date(Date)) %>%
 left_join(features %>% mutate(Date = as.Date(Date)),
 by = c("Store","Date")) %>%
 left_join(stores, by = "Store") %>%
 select(-IsHoliday.y) %>%
 rename(IsHoliday = IsHoliday.x)

``````{r}
missing_tbl <- merged_data %>%
  summarise(
    across(
      .cols = everything(),
      .fns = list(
        count = ~ sum(is.na(.x)),
        percent = ~ mean(is.na(.x)) * 100
      ),
      .names = "{.col}_{.fn}"
    )
  )
```

```
) %>%  
# ----- reshape -----  
  
pivot_longer(  
  cols    = everything(),  
  names_to = c("variable", "metric"),  
  names_sep = "_(?=[^_]+$)", # split on *last* underscore  
  values_to = "value"  
) %>%  
pivot_wider(names_from = metric, values_from = value) %>%  
arrange(desc(percent))  
  
print(missing_tbl)
```

```

## #Handle Missing Values

```
```{r}  
# Remove records with missing sales  
merged_data <- merged_data %>% filter(!is.na(Weekly_Sales))
```

```
# Median imputation for continuous vars  
merged_data <- merged_data %>%  
  mutate(across(starts_with("MarkDown"), ~ replace_na(., 0)))
```

```

```
```{r}  
colSums(is.na(merged_data))
```

```

## #Date & Key Features

```
```{r}  
merged_data <- merged_data %>%  
  mutate(  
    year = year(Date),  
    quarter = quarter(Date),  
    week = isoweek(Date),  
    month = month(Date, label = TRUE),  
    wday_lbl = wday(Date, label = TRUE),  
    ym = yearmonth(Date)  
)
```

```

```
```{r}
```

```
str(merged_data)
```

```

## # Sales Trends

```
#total weekly line plot
```

```
```{r}
```

```

total_sales <- merged_data %>%
  group_by(Date) %>%
  summarise(Total_Weekly_Sales = sum(Weekly_Sales))

ggplot(total_sales, aes(x = Date, y = Total_Weekly_Sales)) +
  geom_line(color = "#2c3e50") +
  labs(title = "Total Weekly Sales Over Time",
       x = "Date", y = "Sales") +
  scale_y_continuous(labels = scales::dollar) +
  theme_minimal()
```
#Weekly sales by store type

```
{r}
merged_data %>%
  group_by(Type) %>%
  summarise(Average_Sales = mean(Weekly_Sales)) %>%
  ggplot(aes(x = Type, y = Average_Sales, fill = Type)) +
  geom_col() +
  labs(title = "Average Weekly Sales by Store Type", x = "Store Type", y = "Average Sales") +
  scale_y_continuous(labels = scales::dollar) +
  theme_minimal()
```
#correlation Hetamap

```
{r}

```

```
library(reshape2)

numeric_vars <- merged_data %>%
  select_if(is.numeric) %>%
  drop_na()

cor_matrix <- cor(numeric_vars)

melted_cormat <- melt(cor_matrix)

ggplot(melted_cormat, aes(Var1, Var2, fill = value)) +
  geom_tile(color = "white") +
  scale_fill_gradient2(low = "red", high = "blue", mid = "white",
    midpoint = 0, limit = c(-1, 1)) +
  theme_minimal() +
  labs(title = "Correlation Heatmap")+
  theme(
    axis.text.x = element_text(angle = 45, hjust = 1)
  )

```
```

```
```

# Correlations as numbers (no plots)
```

```

library(tidyverse)
library(Hmisc)    # rcorr() gives r, p, and n in one shot
library(knitr)    # nice table printing

# 1) Select numeric columns
num_df <- merged_data %>%
  dplyr::select(where(is.numeric))    # same as select_if(is.numeric) but newer

# 2) Correlation matrix with p-values (pairwise complete obs)
rc <- Hmisc::rcorr(as.matrix(num_df), type = "pearson") # use type="spearman" if you prefer
cor_mat <- rc$r # correlations
p_mat <- rc$P # p-values
n_mat <- rc$n # pair counts

# 3) Print the full correlation matrix (rounded)
print(round(cor_mat, 3))

# 4) Tidy, de-duplicated pair list, sorted by |correlation|
cor_pairs <- cor_mat %>%
  as.data.frame() %>%
  rownames_to_column("var1") %>%
  pivot_longer(-var1, names_to = "var2", values_to = "corr") %>%
  # keep each pair once and drop self-correlations
  filter(var1 < var2) %>%
  mutate(
    p = map2_dbl(var1, var2, ~ p_mat[, .y]),

```

```

n = map2_int(var1, var2, ~ n_mat[, .y])
) %>%
arrange(desc(abs(corr)))

# 5) Show the top N strongest correlations
cor_pairs %>%
mutate(corr = round(corr, 3),
p = signif(p, 3)) %>%
slice_head(n = 30) %>%
kable(caption = "Top 30 absolute correlations (Pearson)")

# 6) (Optional) Keep only “meaningful” correlations, e.g.,  $|r| \geq 0.3$  and  $p < 0.05$ 
meaningful_corrs <- cor_pairs %>%
filter(abs(corr) >= 0.30, p < 0.05)

````{r}
merged_data %>%
filter(Weekly_Sales <= 100000) %>%
ggplot(aes(x = Weekly_Sales)) +
geom_histogram(bins = 100, fill = "#3498db", color = "white") +
scale_x_continuous(labels = scales::dollar) +
scale_y_continuous(labels = scales::comma) +
labs(

```

```
title = "Distribution of Weekly Sales (up to $100K)",
x = "Weekly Sales", y = "Count"
) +
theme_minimal()
```

```

#Weekly sales by store type

```
```{r}
```

```
library(tidyverse)
library(lubridate)
library(scales)
```

# 1 — WEEKLY totals by Date × Type × Year

---

```
weekly_type_year <- merged_data %>%
 mutate(Year = factor(year(Date))) %>% # factor → legend / colour
 group_by(Date, Type, Year) %>%
 summarise(weekly_sales = sum(Weekly_Sales), .groups = "drop")
```

# 2 — Vibrant colour palette for the three years

---

```
year_pal <- c("2010" = "#FF5733", # bright orange-red
 "2011" = "#33C1FF", # electric cyan
 "2012" = "#F820FF") # neon magenta
```

# 3 — Plot: raw weekly spikes, coloured by Year, facets by Type

---

```

ggplot(weekly_type_year,
 aes(x = Date, y = weekly_sales,
 colour = Year, group = Year)) +
 geom_line(linewidth = 0.5) + # thin = spikes visible
 facet_wrap(~ Type, ncol = 1, scales = "free_y") +
 scale_colour_manual(
 values = year_pal, name = "Year",
 guide = guide_legend(override.aes = list(linewidth = 2)))
) +
 scale_y_continuous(labels = dollar_format(scale = 1)) +
 labs(
 title = "Weekly Sales Spikes by Store Type (2010–2012)",
 x = NULL, y = "Weekly Sales ($)"
) +
 theme_minimal(base_size = 13) +
 theme(
 axis.text.x = element_text(angle = 45, hjust = 1),
 strip.text = element_text(size = 12, face = "bold"),
 legend.position = "right"
)

```

```

#Seasonality Patterns

Q. Does total sales differ between months, Are there seasonal differences by time period?

```

```{r}

library(dplyr); library(lubridate); library(ggplot2)
library(scales); library(knitr); library(effectsize)

wmt_fy <- merged_data %>%
 mutate(
 fy = if_else(month(Date) >= 2, year(Date), year(Date) - 1L), # fiscal year
 fm = if_else(month(Date) >= 2, month(Date) - 1L, 12L), # fiscal month 1-12
 woy = isoweek(Date) # ISO week (1-52/53)
) %>%
 filter(fy %in% 2010:2011) # keep two complete FYs

```

```

#Monthly profile (Feb → Jan pooled over FY-10 & FY-11)

```

```{r}

— totals per fiscal month (pooled) -----
month_tbl <- wmt_fy %>%
 group_by(fm) %>%
 summarise(total_sales = sum(Weekly_Sales), .groups = "drop")

kable(month_tbl,
 col.names = c("Fiscal Month (1 = Feb)", "Total Sales ($"),
 format.args = list(big.mark = "", scientific = FALSE))

```



```

fm = fiscal month (1 = Feb ... 12 = Jan)

— monthly totals by fiscal year -----
stack_tbl <- wmt_fy %>%
 group_by(fy, fm) %>% # two FYs × 12 months
 summarise(total_sales = sum(Weekly_Sales), .groups = "drop") %>%
 mutate(fm_lab = factor(fm, levels = 1:12,
 labels = c("Feb","Mar","Apr","May","Jun",
 "Jul","Aug","Sep","Oct","Nov","Dec","Jan")),
 fy_lab = factor(fy)) # ensure fy is factor

— stacked bar plot -----
ggplot(stack_tbl, aes(fm_lab, total_sales, fill = fy_lab)) +
 geom_col(position = "stack", colour = "white") +
 scale_fill_brewer(palette = "Set1", name = "Fiscal Year") +
 scale_y_continuous(labels = dollar) +
 labs(title = "Stacked fiscal-month sales (FY-2010 & FY-2011)",
 x = "Fiscal month (Feb = 1 ... Jan = 12)",
 y = "Total sales ($)") +
 theme_minimal(base_size = 11)

```
```
#Week-of-Year profile (ISO weeks 1–52 pooled)
` ` ` {r}

```

```
woy_tbl <- wmt_fy %>%
 group_by(woy) %>%
 summarise(total_sales = sum(Weekly_Sales), .groups = "drop")
```

```
kable(woy_tbl,
 col.names = c("ISO Week", "Total Sales ($)'),
 format.args = list(big.mark = ",", scientific = FALSE))
```

```
top3 <- woy_tbl %>% slice_max(total_sales, n = 3)
```

```
bot3 <- woy_tbl %>% slice_min(total_sales, n = 3)
```

```
w_gap_abs <- max(top3$total_sales) - min(bot3$total_sales)
```

```
w_gap_pct <- 100 * w_gap_abs / max(top3$total_sales)
```

1

```{r}

— weekly totals pooled over FY-10 & 11 —————

```
woy_tot <- wmt_fy %>%
```

```
group_by(woy) %>% # iso week number (1-52)
```

```
summarise(total_sales = sum(Weekly_Sales), .groups = "drop")
```

— line plot -----

```
ggplot(woy_tot, aes(woy, total_sales)) +
```

```
geom_line(colour = "#2980b9", linewidth = 0.8) +  
  geom_point(size = 1.5, colour = "#2980b9") +  
  scale_x_continuous(breaks = seq(1, 52, by = 4)) +  
  scale_y_continuous(labels = dollar) +  
  labs(title = "ISO week-of-year sales (FY-2010 & FY-2011 pooled)",  
       x = "ISO week number",  
       y = "Total sales ($)") +  
  theme_minimal(base_size = 11)
```

```

---

## STORE FISCAL ANALYSIS

Which stores are performing best/worst and How do stores compare systematically?

```
```{r}  
library(dplyr)  
library(lubridate)  
library(tidyr)  
library(ggplot2)  
library(viridis)  
library(scales)  
library(knitr)
```

```
# — 0 · fiscal-year helper -----
```

```
merged_fy <- merged_data %>%  
  mutate(
```

```

fy = if_else(month(Date) >= 2, year(Date), year(Date) - 1L), # FY key

fy = factor(fy, levels = 2010:2012, labels = c("FY2010","FY2011","FY2012"))

)

# — 1.1 Year-by-store aggregates -----
store_year_tbl <- merged_fy %>%
  group_by(Store, fy) %>%
  summarise(year_sales = sum(Weekly_Sales, na.rm = TRUE), .groups = "drop")

store_wide <- store_year_tbl %>%
  pivot_wider(names_from = fy,
  values_from = year_sales,
  names_prefix = "Sales_") %>%
  mutate(Sales_Total = Sales_FY2010 + Sales_FY2011 + Sales_FY2012)

# — 1.2 Rank & category flags -----
ranked_tbl <- store_wide %>%
  arrange(desc(Sales_Total)) %>%
  mutate(
    Rank_Overall = row_number(),
    Rank_FY2012 = rank(-Sales_FY2012, ties.method = "min"),
    Category = case_when(
      Rank_Overall <= 3 ~ "Top 3",
      Rank_Overall > n() - 3 ~ "Bottom 3",
      TRUE ~ "Middle"
    )
  )

```

```

kable(ranked_tbl,
      col.names = c("Store","Sales FY10","Sales FY11","Sales FY12*",
                  "Sales Total","Rank Overall","Rank FY12","Category"),
      format.args = list(big.mark = "", scientific = FALSE),
      caption = "*FY-2012 covers Feb → Oct only")

# — 2 · heat-map -----
heat_df <- ranked_tbl %>%
  select(Store, Sales_FY2010, Sales_FY2011, Sales_FY2012, Sales_Total, Category) %>%
  pivot_longer(starts_with("Sales_"),
               names_to = "Year",
               values_to = "Sales") %>%
  mutate(Store = factor(Store, levels = ranked_tbl$Store)) # best → worst

ggplot(heat_df, aes(Year, Store, fill = Sales)) +
  geom_tile(colour = "white", linewidth = 0.3) +
  scale_fill_viridis(option = "C", labels = dollar,
                     name = "Sales $") +
  # red outline for Top / Bottom stores
  geom_tile(data = heat_df %>% filter(Category != "Middle"),
            colour = "red", linewidth = 0.8, fill = NA) +
  labs(title = "Store performance heat-map by fiscal year (FY-2010 – FY-2012*)",
       subtitle = "*FY-2012 bar shorter because Nov–Jan withheld",
       x = NULL, y = "Store (ranked best → worst)") +
  theme_minimal(base_size = 10) +
  theme(axis.text.y = element_text(size = 6))

```

```

---

```
#DEPARTMENT
```

```
#Department-Level Performance by Fiscal Year
```

```
#Q. What is the sales distribution across departments and how stable is this
distribution
```

```
over time?
```

```
```{r}
```

```
library(dplyr); library(lubridate); library(tidyr)  
library(ggplot2); library(scales); library(knitr)
```

```
dept_year_tbl <- merged_data %>%  
  mutate(fy = if_else(month(Date) >= 2, year(Date), year(Date) - 1L),  
    fy = factor(fy, levels = 2010:2012,  
    labels = c("FY2010","FY2011","FY2012")) %>%  
  group_by(Dept, fy) %>%  
  summarise(year_sales = sum(Weekly_Sales), .groups = "drop")
```

```
dept_wide <- dept_year_tbl %>%  
  pivot_wider(names_from = fy, values_from = year_sales,  
  names_prefix = "Sales_") %>%  
  mutate(Sales_Total = Sales_FY2010 + Sales_FY2011 + Sales_FY2012) %>%  
  arrange(desc(Sales_Total)) %>%  
  mutate(Rank_Overall = row_number())
```

```
# absolute & % gap between best and worst
```

```

abs_gap <- dept_wide$Sales_Total[1] - dept_wide$Sales_Total[nrow(dept_wide)]
pct_gap <- 100 * abs_gap / dept_wide$Sales_Total[1]

kable(dept_wide,
      col.names = c("Dept","Sales FY10","Sales FY11","Sales FY12*",
                  "Sales Total","Rank Overall"),
      format.args = list(big.mark = "", scientific = FALSE),
      caption = paste0("*FY-2012 covers Feb → Oct only | Gap best→worst = $",
                      comma(abs_gap), " (", round(pct_gap,1), "%)"))

````

#Top-3-department share plot
```{R}

# identify top-3 depts by pooled sales
top3 <- dept_wide$Dept[1:3]

plot_df <- merged_data %>%
  filter(Dept %in% top3) %>%
  mutate(fy = if_else(month(Date) >= 2, year(Date), year(Date)-1L),
         fy = factor(fy, levels = 2010:2012,
                     labels = c("FY2010","FY2011","FY2012")))) %>%
  group_by(fy) %>%
  summarise(top3_sales = sum(Weekly_Sales), .groups = "drop") %>%
  # chain totals for denominator
  left_join(
    merged_data %>%
      mutate(fy = if_else(month(Date) >= 2, year(Date), year(Date)-1L),

```

```

fy = factor(fy, levels = 2010:2012,
            labels = c("FY2010","FY2011","FY2012")) %>%
group_by(fy) %>% summarise(chain_sales = sum(Weekly_Sales), .groups = "drop"),
by = "fy") %>%
mutate(pct_share = top3_sales / chain_sales)

ggplot(plot_df, aes(fy, pct_share)) +
  geom_col(fill = "#FF5733") +
  geom_text(
    aes(label = scales::percent(pct_share, accuracy = 1)),
    vjust = -0.3,           # just above each bar
    size = 4,               # text size in mm
    fontface = "bold"
  ) +
  scale_y_continuous(labels = percent_format(accuracy = 1),
                     limits = c(0, 0.50), expand = c(0,0)) +
  labs(title = "Share of Chain Sales from Top-3 Departments",
       subtitle = "Fiscal-year view (FY-2012 partial Feb–Oct)",
       x = NULL, y = "Percent of Total Sales") +
  theme_minimal(base_size = 11)

```
```
```{r}
top3 <- dept_wide$Dept[1:3]

```

```

--- Totals by fiscal year (chain & top-3) ---
chain_by_fy <- dept_year_tbl %>%
 group_by(fy) %>%
 summarise(chain_sales = sum(year_sales), .groups = "drop")

top3_by_fy <- dept_year_tbl %>%
 filter(Dept %in% top3) %>%
 group_by(fy) %>%
 summarise(top3_sales = sum(year_sales), .groups = "drop")

--- Each non-top-3 department's share of the "rest 80%" (per FY) ---
rest_dept_fy <- dept_year_tbl %>%
 filter(!(Dept %in% top3)) %>%
 left_join(chain_by_fy, by = "fy") %>%
 left_join(top3_by_fy, by = "fy") %>%
 mutate(
 rest_pool = chain_sales - top3_sales,
 share_of_rest_fy = if_else(rest_pool > 0, year_sales / rest_pool, NA_real_),
 share_of_total_fy = year_sales / chain_sales
)

--- Average contributions across FYs ---
rest_avg <- rest_dept_fy %>%
 group_by(Dept) %>%
 summarise(
 pooled_share_of_rest = sum(year_sales, na.rm = TRUE) / sum(rest_pool, na.rm =
TRUE),

```

```

mean_share_of_rest = mean(share_of_rest_fy, na.rm = TRUE),
pooled_share_of_total = sum(year_sales, na.rm = TRUE) / sum(chain_sales, na.rm =
TRUE),
.groups = "drop"
) %>%
arrange(desc(pooled_share_of_rest))

sum_rest_share <- sum(rest_avg$pooled_share_of_rest, na.rm = TRUE)
message("Sum of pooled shares across rest = ", round(sum_rest_share, 3))

rest_tbl <- rest_avg %>%
mutate(
`Pooled share of rest` = scales::percent(pooled_share_of_rest, accuracy = 0.1),
`Mean share of rest (FY avg)` = scales::percent(mean_share_of_rest, accuracy = 0.1),
`Pooled share of TOTAL` = scales::percent(pooled_share_of_total, accuracy = 0.1)
) %>%
select(Dept, `Pooled share of rest`, `Mean share of rest (FY avg)`, `Pooled share of
TOTAL`)

knitr::kable(rest_tbl %>% slice_head(n = 15),
caption = "Top 15 contributors among the OTHER ~80% (excl. top-3 depts)")

```

```

```
```{r}
```

```
library(forcats); library(scales); library(ggplot2); library(dplyr)
```

```
N <- 15
```

```
rest_top <- rest_avg %>%
 arrange(desc(pooled_share_of_rest)) %>%
 mutate(Dept = as.factor(Dept))
```

```
rest_topN <- rest_top %>% slice_head(n = N)
others_share <- 1 - sum(rest_topN$pooled_share_of_rest, na.rm = TRUE)
```

```
rest_topN_plus <- bind_rows(
 rest_topN,
 tibble(Dept = factor("Other depts"), pooled_share_of_rest = others_share)
) %>%
 mutate(Dept = fct_reorder(Dept, pooled_share_of_rest))
```

```
ggplot(rest_topN_plus, aes(x = Dept, y = pooled_share_of_rest)) +
 geom_col(fill = "#2c3e50") +
 coord_flip() +
 geom_text(aes(label = percent(pooled_share_of_rest, accuracy = 0.1)),
 hjust = -0.1, size = 3) +
 scale_y_continuous(labels = percent, expand = expansion(mult = c(0, .15))) +
 labs(
 title = paste0("Which depts make up the OTHER ~80%? (Top ", N, " + 'Other')"),
 x = "Department (non-top-3 only)",
```

y = "Share of the remaining ~80% (pooled across FYs)"

) +

```
theme_minimal(base_size = 11)
```

10

```{r}

```
library(treemapify)
```

```
ggplot(rest_top, aes(area = pooled_share_of_rest, fill = pooled_share_of_rest,
```

label = Dept)) +

```
geom_treemap() +
```

```
geom_treemap_text(colour = "white", place = "centre", grow = TRUE, reflow = TRUE) +
```

```
scale_fill_gradient(low = "#9ecae1", high = "#08519c", labels = percent) +
```

labs(

title = "Treemap: contribution of each non-top-3 department to the OTHER ~80%",

fill = "Share of rest"

) +

```
theme_minimal(base_size = 11) + theme(legend.position = "right")
```

Fiscal-Year Monthly-Totals Analysis

Year over year

#.Are there statistically significant differences in sales performance across fiscal years
2010-2012?

```
```{r}
```

```
library(dplyr); library(lubridate); library(tidyr)
library(ggplot2); library(scales); library(knitr)
library(effectsize); library(broom)
```

```
```
```

Aggregate monthly totals by fiscal year (Feb → Oct only)

```
```{r}
```

```
fy_month_tbl <- merged_data %>%
---- fiscal keys -----
mutate(
 fy = if_else(month(Date) >= 2, year(Date), year(Date) - 1L),
 fy = factor(fy, levels = 2010:2012, labels = c("FY2010","FY2011","FY2012")),
 Month = month(Date, label = TRUE, abbr = TRUE)
) %>%
---- keep months present in all 3 FYs (Feb–Oct) -----
filter(Month %in% month(2:10, label = TRUE, abbr = TRUE)) %>%
group_by(fy, Month) %>%
summarise(total_sales = sum(Weekly_Sales), .groups = "drop")
```

```
```
```

```
```{r}
```

```
fy_month_wide <- fy_month_tbl %>%
pivot_wider(names_from = fy, values_from = total_sales)
```

```
kable(fy_month_wide,
col.names = c("Month","FY-2010","FY-2011","FY-2012*"),
format.args = list(big.mark = "", scientific = FALSE),
caption = "*FY-2012 covers Feb → Oct only")
```

```
```
```

Grouped-bar plot

```
```{r}
```

```
ggplot(fy_month_tbl, aes(Month, total_sales, fill = fy)) +
geom_col(position = "dodge") +
scale_y_continuous(labels = dollar) +
scale_fill_brewer(palette = "Set2", name = "Fiscal Year") +
labs(title = "Monthly sales by fiscal year (common window Feb–Oct)",
x = NULL, y = "Total sales ($)") +
theme_minimal(base_size = 11)
```

```
```
```

```

```{r}

library(dplyr); library(lubridate); library(broom)

library(effectsize); library(knitr)

--- weekly chain totals, Feb–Oct window -----
weekly_common <- merged_data %>%
 mutate(
 fy = if_else(month(Date) >= 2, year(Date), year(Date)-1L),
 fy = factor(fy, levels = 2010:2012,
 labels = c("FY2010","FY2011","FY2012")),
 Month = month(Date, label = TRUE, abbr = TRUE)
) %>%
 filter(Month %in% month(2:10, label = TRUE, abbr = TRUE)) %>%
 group_by(fy, Date) %>%
 summarise(chain_sales = sum(Weekly_Sales), .groups = "drop")

--- choose parametric vs non-parametric -----
aov_mod <- aov(chain_sales ~ fy, data = weekly_common)
norm_p <- shapiro.test(residuals(aov_mod))$p.value

if (norm_p > .05) {
 test_name <- "One-way ANOVA"
 stat_out <- tidy(aov_mod)[1, c("df","statistic","p.value")]
 eta_tbl <- effectsize::eta_squared(aov_mod, partial = TRUE)
 term_col <- intersect(c("Effect","Parameter","Term"), names(eta_tbl))[1]
 eff_sz <- eta_tbl %>%
 filter(.data[[term_col]] == "fy") %>%
 select(matches("^Eta2")) %>% pull() %>% round(3)
}

```

```

} else {

 test_name <- "Kruskal-Wallis"

 kw_mod <- kruskal.test(chain_sales ~ fy, data = weekly_common)

 stat_out <- tidy(kw_mod)[, c("parameter","statistic","p.value")] %>%
 rename(df = parameter)

 eff_sz <- effectsize::epsilon_squared(kw_mod)$Epsilon2 %>% round(3)

}

cat("\n***", test_name, "***\n\n") # bold header

kable(stat_out, digits = 3,
 col.names = c("df", ifelse(test_name=="One-way ANOVA","F","H"), "p"))

cat("\nPartial η^2 (fy) =", eff_sz, "\n")

````

#which year is higher/lower
```{r}
weekly_common %>%
 group_by(fy) %>%
 summarise(Mean = mean(chain_sales),
 SD = sd(chain_sales),
 n = n()) %>%
 mutate(across(Mean:SD, scales::dollar)) %>%
 knitr::kable()

```

```

#What is the quantitative impact of promotional activities and holiday periods on weekly sales performance?

#What is the effect of promotions/holidays

Fiscal-Years 2010 & 2011 — Promotion- and Holiday-Week Impact

```{R}

```
library(dplyr); library(lubridate); library(ggplot2)
library(scales); library(effsize); library(knitr)
```

# — Filter to the two complete fiscal years -----

```
fy_span <- merged_data %>%
filter(Date >= as.Date("2010-02-05"),
 Date <= as.Date("2012-01-27"))
```

# — Promo flag (any Markdown > 0) -----

```
fy_span <- fy_span %>%
mutate(promo_flag = as.integer(if_any(starts_with("MarkDown"), ~ .x > 0)))
```

# — Weekly aggregation (one row per Friday) -----

```
weekly_df <- fy_span %>%
group_by(Date) %>%
summarise(total_sales = sum(Weekly_Sales, na.rm = TRUE),
```

```

promo_flag = first(promo_flag),
holiday = first(IsHoliday), .groups = "drop")

````

````{r}

---- non-parametric test & effect size (PROMO) -----
Was: t.test(...) + Cohen's d

w_promo <- wilcox.test(total_sales ~ promo_flag,
 data = weekly_df,
 exact = FALSE, conf.int = TRUE, alternative = "two.sided")

delta_promo <- effsize::cliff.delta(total_sales ~ promo_flag,
 data = weekly_df)$estimate

````

````{r}

promo_tbl <- tibble(
 Group = c("No-Promo (0)", "Promo (1)", " Δ (1-0)"),
 Mean_Weekly_Sales = c(promo_stats$Mean, diff(promo_stats$Mean)),
 SD = c(promo_stats$SD, NA),
 n = c(promo_stats$n, NA),
 `95% CI` = c("", "", paste0(scales::dollar(w_promo$conf.int[1]),
 " → ", scales::dollar(w_promo$conf.int[2]))),
 `p-value` = c("", "", signif(w_promo$p.value, 3)),
 `Cliff's delta` = c("", "", round(delta_promo, 2))
)

```

```

knitr::kable(promo_tbl, digits = 0,
 col.names = c("Group","Mean ($)","SD","n",
 "95 % CI","p-value","Cliff's Δ"))

```
```

```{r}
# ---- non-parametric test & effect size (HOLIDAY) -----
w_hol<- wilcox.test(total_sales ~ holiday,
                      data = weekly_df,
                      exact = FALSE, conf.int = TRUE, alternative = "two.sided")
delta_hol<- effsize::cliff.delta(total_sales ~ holiday,
                                   data = weekly_df)$estimate

hol_tbl<- tibble(
  Group      = c("Non-Holiday (0)", "Holiday (1)", "Δ (1-0)"),
  Mean_Weekly_Sales = c(hol_stats$Mean, diff(hol_stats$Mean)),
  SD         = c(hol_stats$SD, NA),
  n          = c(hol_stats$n, NA),
  `95% CI`   = c("", "", paste0(scales::dollar(w_hol$conf.int[1]),
                                 " → ", scales::dollar(w_hol$conf.int[2]))),
  `p-value`  = c("", "", signif(w_hol$p.value, 3)),
  `Cliff's delta` = c("", "", round(delta_hol, 2)))

```

```
)
```

```
knitr::kable(hol_tbl, digits = 0,  
             col.names = c("Group","Mean ($)","SD","n",  
                          "95 % CI","p-value","Cliff's Δ"))  
  
` ` ` {r}  
library(ggplot2); library(scales)  
  
weekly_df <- merged_data %>%  
  mutate(promo_flag = as.integer(if_any(starts_with("MarkDown"), ~ .x > 0))) %>%  
  group_by(Date) %>%  
  summarise(total_sales = sum(Weekly_Sales),  
            holiday = max(IsHoliday),  
            promo_flag = max(promo_flag),  
            .groups = "drop")  
  
ggplot(weekly_df, aes(Date, total_sales)) +  
  geom_line(colour = "#34495e", linewidth = 0.8) +  
  geom_point(data = weekly_df %>% filter(promo_flag==1),  
             aes(Date, total_sales), alpha = 0.5, size = 1.2) +  
  geom_vline(data = weekly_df %>% filter(holiday==1),  
             aes(xintercept = Date), colour = "#e74c3c", linetype = "dashed", linewidth = 0.6) +  
  scale_y_continuous(labels = dollar) +  
  labs(title = "Chain weekly sales with promo points and holiday markers",  
       x = NULL, y = "Weekly sales ($)") +  
  theme_minimal(base_size = 11)
```

```

---

```{R}

```
library(tsibble); library(feasts); library(anomalize); library(dplyr)
```

```
chain_ts <- merged_data %>%
```

```
  group_by(Date) %>%
```

```
  summarise(Weekly_Sales = sum(Weekly_Sales), .groups = "drop") %>%
```

```
  as_tsibble(index = Date)
```

```
# STL decomposition
```

```
chain_ts %>%
```

```
  model(STL(Weekly_Sales ~ season(window = "periodic") + trend())) %>%
```

```
  components() %>%
```

```
  autoplot() + ggtitle("STL decomposition (chain weekly)")
```

```

```
```{r}
```

```
library(dplyr)
```

```
library(tsibble)
```

```
library(feasts)
```

```
library(ggplot2)
```

```
library(tseries)
```

```
# 1) Chain-level weekly totals -----
```

```
chain_tbl <- merged_data %>%
```

```
  group_by(Date) %>%
```

```
  summarise(Weekly_Sales = sum(Weekly_Sales, na.rm = TRUE), .groups = "drop") %>%
```

```
  arrange(Date)
```

```
# tsibble for ACF/PACF (ggplot style)
```

```
chain_tsbl <- chain_tbl %>% as_tsibble(index = Date)
```

```
chain_tsbl %>%
```

```
  ACF(Weekly_Sales, lag_max = 60) %>%
```

```
  autoplot() + ggtitle("ACF (Chain Weekly_Sales)")
```

```
chain_tsbl %>%
```

```
  PACF(Weekly_Sales, lag_max = 60) %>%
```

```
  autoplot() + ggtitle("PACF (Chain Weekly_Sales)")
```

```
chain_ts <- ts(chain_tbl$Weekly_Sales, frequency = 52)
```

```
adf_raw <- adf.test(chain_ts, alternative = "stationary")
print(adf_raw)
```

```

---

```
#Feature engineering & Modeling
```

```
```{r}
```

```
# -----
```

```
# Final modeling
```

```
# -----
```

```
# Load libraries
```

```
library(tidyverse)
```

```
library(lubridate)
```

```
library(slider)
```

```
library(forecast)
```

```
library(prophet)
```

```
library(randomForest)
```

```
library(xgboost)
```

```
library(purrr)
```

```
library(ranger)
```

```
library(vip)
```

```

#-----
# 1) Feature engineering
#-----

df_chain <- merged_data %>%
  group_by(Date) %>%
  summarise(Weekly_Sales = sum(Weekly_Sales), .groups = "drop") %>%
  arrange(Date) %>%
  mutate(
    lag_1 = lag(Weekly_Sales, 1),
    lag_4 = lag(Weekly_Sales, 4),
    lag_52 = lag(Weekly_Sales, 52),
    roll_mean_4 = slide_dbl(Weekly_Sales, mean, .before = 3, .complete = TRUE),
    roll_mean_13 = slide_dbl(Weekly_Sales, mean, .before = 12, .complete = TRUE),
    week = isoweek(Date),
    month = month(Date),
    year = year(Date),
    sin52 = sin(2 * pi * week / 52),
    cos52 = cos(2 * pi * week / 52)
  ) %>%
  drop_na()

```

```

#-----
# 2) Monte-Carlo CV setup
#-----


make_time_mc_cv <- function(data, train_size, test_size, n_reps) {
  n <- nrow(data)
  splits <- vector("list", n_reps)
  set.seed(123)

```

```

for (i in seq_len(n_reps)) {
  start <- sample(1:(n - train_size - test_size + 1), 1)
  train_idx <- start:(start + train_size - 1)
  test_idx <- (start + train_size):(start + train_size + test_size - 1)
  splits[[i]] <- list(train = data[train_idx, ], test = data[test_idx, ])
}
splits
}

cv_splits <- make_time_mc_cv(df_chain, train_size = 80, test_size = 8, n_reps = 10)

#-----
# 3) Model fitting function

# (3 classical + 3 ML models)

#-----
run_models <- function(split) {
  train <- split$train
  test <- split$test

  x_vars <- c(
    "lag_1","lag_4","lag_52",
    "roll_mean_4","roll_mean_13",
    "week","month","year","sin52","cos52"
  )
}

# ----- Classical -----
ts_train <- ts(train$Weekly_Sales, frequency = 52)

fit_sarima <- auto.arima(ts_train, seasonal = TRUE)

```

```

fc_sarima <- forecast(fit_sarima, h = nrow(test))

sarima_pred <- as.numeric(fc_sarima$mean)

fit_ets <- ets(ts_train)

fc_ets <- forecast(fit_ets, h = nrow(test))

ets_pred <- as.numeric(fc_ets$mean)

fit_snaive <- forecast::snaive(ts_train, h = nrow(test), lag = 52)

snaive_pred <- as.numeric(fit_snaive$mean)

# ----- Machine learning -----

train_prophet <- train %>% dplyr::select(ds = Date, y = Weekly_Sales)

m_prophet <- prophet(

  train_prophet,

  yearly.seasonality = TRUE,

  weekly.seasonality = FALSE,

  daily.seasonality = FALSE,

  verbose = FALSE

)

future <- test %>% dplyr::select(ds = Date)

fc_prophet <- predict(m_prophet, future)

prophet_pred <- fc_prophet$yhat

rf_fit <- randomForest(

  Weekly_Sales ~ .,

  data = dplyr::select(train, Weekly_Sales, dplyr::all_of(x_vars)),

  ntree = 200, mtry = 3

)

```

```

rf_pred <- predict(rf_fit, newdata = dplyr::select(test, dplyr::all_of(x_vars)))

dtrain <- xgb.DMatrix(as.matrix(train[, x_vars]), label = train$Weekly_Sales)
dtest <- xgb.DMatrix(as.matrix(test[, x_vars]))

xgb_fit <- xgboost(
  data = dtrain,
  nrounds = 200,
  objective = "reg:squarederror",
  max_depth = 4,
  eta = 0.1,
  subsample = 0.8,
  colsample_bytree = 0.8,
  verbose = 0
)
xgb_pred <- predict(xgb_fit, dtest)

# Return long format
tibble::tibble(
  Date = test$Date,
  actual = test$Weekly_Sales,
  SARIMA = sarima_pred,
  ETS = ets_pred,
  sNaive = snaive_pred,
  Prophet = prophet_pred,
  RandomForest = rf_pred,
  XGBoost = xgb_pred
) %>%
  tidyverse::pivot_longer(

```

```

cols = -c(Date, actual),
names_to = "model",
values_to = "pred"

)

}

# Run CV

results <- purrr::map_dfr(cv_splits, run_models, .progress = TRUE)

#-----
# 4) Evaluate model performance
#-----

error_metrics <- results %>%
  filter(!is.na(actual), !is.na(pred)) %>%
  group_by(model) %>%
  summarise(
    MAE = mean(abs(actual - pred)),
    RMSE = sqrt(mean((actual - pred)^2)),
    MAPE = mean(abs(actual - pred) / actual) * 100,
    .groups = 'drop'
  ) %>%
  arrange(MAPE)
print(error_metrics)

#-----
# 5) Diagnostic plots
#-----

# Boxplot of absolute errors by model

```

```

ggplot(results, aes(x = model, y = abs(actual - pred), fill = model)) +
  geom_boxplot() +
  labs(title = "Absolute Error Distribution by Model", x = NULL, y = "Absolute Error") +
  theme_minimal()

# Actual vs. predicted (averaged over folds)
avg_preds <- results %>%
  group_by(Date, model) %>%
  summarise(actual = mean(actual), pred = mean(pred), .groups = "drop")

ggplot(avg_preds, aes(x = Date)) +
  geom_line(aes(y = actual, colour = "Actual"), size = 0.8) +
  geom_line(aes(y = pred, colour = model), linetype = "solid") +
  labs(
    title = "Actual vs Predicted Weekly Sales (averaged over CV folds)",
    x = "Date", y = "Weekly Sales", colour = "Series"
  ) +
  theme_minimal()

#-----
# 6) Feature importance (full data)
#-----

# Random Forest (ranger) importance
rf_full <- ranger(
  Weekly_Sales ~ .,
  data = dplyr::select(
    df_chain, Weekly_Sales,
    dplyr::all_of(c(

```

```
"lag_1","lag_4","lag_52","roll_mean_4","roll_mean_13",
"week","month","year","sin52","cos52"
)))
),
importance = "permutation",
num.trees = 500,
seed = 1
)
vip::vip(rf_full, num_features = 10)

# XGBoost importance
dall <- xgb.DMatrix(
  data = as.matrix(df_chain[, c(
    "lag_1","lag_4","lag_52","roll_mean_4","roll_mean_13",
    "week","month","year","sin52","cos52"
  )]),
  label = df_chain$Weekly_Sales
)
xgb_full <- xgboost(
  data = dall,
  nrounds = 300,
  objective = "reg:squarederror",
  max_depth = 4,
  eta = 0.1,
  subsample = 0.8,
  colsample_bytree = 0.8,
  verbose = 0
)
```

```
xgb_imp <- xgb.importance(model = xgb_full)

ggplot(xgb_imp, aes(x = reorder(Feature, Gain), y = Gain)) +
  geom_col(fill = "#e74c3c") +
  coord_flip() +
  labs(title = "XGBoost – Feature Importance (Gain)",
       x = "Feature", y = "Gain") +
  theme_minimal()
```

```
ggplot(avg_preds, aes(x = Date)) +
  geom_line(aes(y = actual, colour = "Actual"), size = 0.8) +
  geom_line(aes(y = pred, colour = "Predicted"), linetype = "dashed") +
  facet_wrap(~ model, scales = "free_y", ncol = 2) +
  labs(
    title = "Test set: Actual vs Predicted Overlay by Model",
    y = "Weekly Sales", colour = ""
  ) +
  theme_minimal()
```

```

```
#split by split results
```

```
```{r warning=FALSE}

# Apply run_models to each split and add an identifier for the fold
```

```
results <- purrr::imap_dfr(  
  cv_splits,  
  ~ run_models(.x) %>% mutate(split_id = .y),  
  .progress = TRUE  
)  
  
# Now compute metrics for each split and model  
metrics_by_split <- results %>%  
  filter(!is.na(actual), !is.na(pred)) %>%  
  group_by(split_id, model) %>%  
  summarise(  
    MAE = mean(abs(actual - pred)),  
    RMSE = sqrt(mean((actual - pred)^2)),  
    MAPE = mean(abs(actual - pred) / actual) * 100,  
    .groups = 'drop'  
) %>%  
  arrange(split_id, MAPE)  
  
print(metrics_by_split)  
  
```{r}  

7. Classical residual diagnostics (full series)
```

```

library(forecast)

Fit SARIMA, ETS and Prophet on the full chain-level series
ts_full <- ts(df_chain$Weekly_Sales, frequency = 52)

SARIMA on full series
fit_sarima_full <- auto.arima(ts_full, seasonal = TRUE)
sarima_resid_full <- residuals(fit_sarima_full)

ETS on full series
fit_ets_full <- ets(ts_full)
ets_resid_full <- residuals(fit_ets_full)

---- Prophet with 52-week seasonality + holidays + exogenous ----
df_prophet_full <- df_chain %>%
 transmute(
 ds = Date, y = Weekly_Sales,
 week = isoweek(Date),
 month = month(Date),
 sin52 = sin(2*pi*week/52),
 cos52 = cos(2*pi*week/52)
)

m_prophet_full <- prophet(
 yearly.seasonality = FALSE, # we'll add a custom 52-week seasonality
 weekly.seasonality = FALSE, # not meaningful for weekly data
```

```

daily.seasonality = FALSE,
seasonality.mode = "multiplicative",
changepoint.prior.scale = 0.2,
holidays.prior.scale = 10,
seasonality.prior.scale = 10,
verbose = FALSE
)

custom annual cycle (52 weeks)
m_prophet_full <- add_seasonality(
 m_prophet_full, name = "annual52", period = 52, fourier.order = 10
)

(optional) add US holidays
m_prophet_full <- add_country_holidays(m_prophet_full, country_name = "US")

add exogenous regressors (Fourier proxies, calendar)
for (rv in c("sin52","cos52","week","month")){
 m_prophet_full <- add_regressor(m_prophet_full, rv)
}

m_prophet_full <- fit.prophet(m_prophet_full, df_prophet_full)

p_full <- predict(m_prophet_full, df_prophet_full)
prophet_resid_full <- df_prophet_full$y - p_full$yhat

Residual diagnostics
forecast::ggAcf(prophet_resid_full) + labs(title = "Prophet Residuals – ACF")

```

```
forecast::ggPacf(prophet_resid_full) + labs(title = "Prophet Residuals – PACF")
Box.test(prophet_resid_full, lag = 24, type = "Ljung-Box")

Plot residual ACF and PACF for each classical model
These lines produce separate ACF/PACF plots; you can display them as needed
SARIMA residual ACF/PACF
ggAcf(sarima_resid_full) + labs(title = "SARIMA Residuals – ACF")
ggPacf(sarima_resid_full) + labs(title = "SARIMA Residuals – PACF")
Ljung–Box test
Box.test(sarima_resid_full, lag = 24, type = "Ljung-Box")

ETS residual ACF/PACF
ggAcf(ets_resid_full) + labs(title = "ETS Residuals – ACF")
ggPacf(ets_resid_full) + labs(title = "ETS Residuals – PACF")
Box.test(ets_resid_full, lag = 24, type = "Ljung-Box")
```

```